

# KratorAI Brand Onboarding — Product Implementation Brief

## What This Document Is

This is a product specification for the KratorAI brand knowledge extraction onboarding flow. It covers the full experience design, the question architecture, the voice/type input system, and how the captured data gets structured into a Brand Intelligence Profile. Developer and designer should read this together — the decisions in each discipline directly affect the other.

---

## The Core Design Principle

Before anything else, both the developer and designer need to internalize one principle that governs every decision in this document:

**This is not a form. It is a conversation.**

Every time a design or engineering decision makes this feel like a form — a field to fill in, a required input, a progress bar counting steps — it is the wrong decision. Every time a decision makes it feel like a conversation — a question being asked, a response being heard, an intelligent follow-up — it is the right one.

The business owner should leave onboarding feeling like they just talked to someone who genuinely understood their brand. Not like they just completed a setup wizard.

---

## Overview of What the Onboarding Must Accomplish

The onboarding has one job: build a Brand Intelligence Profile that is rich enough for KratorAI's AI to generate content that sounds like this specific business and no other.

That profile has five layers, each progressively harder to extract:

1. **What they sell** — products, services, prices. Easy. Most onboarding flows get this.
2. **Who they sell to** — audience description, customer psychology. Moderate. Most flows get a surface version of this.
3. **How they speak** — tone, vocabulary, rhythm, cultural register. Hard. Most flows fail here because they ask people to describe their voice instead of demonstrating it.
4. **What they actually believe** — their genuine opinions, their point of view, what makes them distinct as a person not just as a business. Very hard. Almost no tool captures this at all.
5. **How they naturally express themselves** — not described, but captured live through their actual words, spoken or written without self-editing.

The onboarding must reach layer five for every user. The design decisions below are all in service of that goal.

---

## The Full Experience Arc — What the User Should Feel at Each Stage

Design every screen against this emotional sequence. If a screen doesn't serve one of these emotional states, reconsider why it exists.

**Anticipation** — Something different is about to happen here.

**Engagement** — This is actually a conversation, not a form.

**Surprise** — It already sounds like me and we're not even done.

**Confidence** — It genuinely understood my brand.

**Excitement** — Thirty days of content, ready, looking like this.

**Relief** — I don't have to do anything now.

**Awe** — It just posted for me while I was doing something else.

The onboarding covers the first five. The last two happen post-onboarding and need to be designed for separately — but the onboarding sets them up.

---

# Part One: The Opening Screen

## What it must do

Set the frame. Signal that this is different from every other onboarding the user has experienced. Create anticipation without overwhelming.

## What it must not do

Start with a progress bar. List the number of steps. Use the word "setup." Show a wall of fields.

## The copy approach

The opening screen should carry a confident, human statement of what is about to happen. Something in this direction — the exact copy should be refined by your writer, but the content should be:

- Specific about what KratorAI will learn
- Clear that the result happens at the end
- Honest that it takes real input to work well
- Warm but not salesy

An example of the right tone: *"In the next few minutes, KratorAI is going to learn everything it needs to know about your brand — how you speak, who you sell to, what makes you different. By the time we're done, it'll sound like you."*

No bullet points listing features. No checkbox list of what onboarding covers. Just that statement and a single call to action to begin.

## The very first question

Before entering the main question flow, ask one simple question in plain human language: the user's name and their business name. This is the first moment the voice mode is active — but the stakes are zero. They're just saying their name. This is deliberate. It introduces the voice interface at the lowest possible anxiety moment so it feels normal before anything demanding is asked.

---

# Part Two: The Input System — Voice and Type

## The fundamental decision

Both voice and type are available throughout the onboarding. They are not equal choices presented at every question. The interface has a **designed default for each question** — voice for some, type for others — with a quiet, always-available toggle to switch.

The designer's job is to make the default feel like the natural choice for that question, so the toggle is rarely needed but always available.

## Voice default questions — these questions extract personality and natural expression

The following questions should default to voice input because the signal lives in *how* the person speaks, not in what they carefully type:

- Describe your business in plain language
- Who is your best customer — describe them as a person
- What do you believe that competitors wouldn't say publicly
- What do your customers most often misunderstand
- What are you most proud of about your business
- If your brand was a person at a dinner party, how would they behave

For these questions, the microphone should be **already listening when the question appears**. There is no button to press to begin speaking. The interface detects voice activity the same way a conversation does — you just start talking. There is a subtle visual indicator that the system is listening, but it should not resemble a recording button. It should resemble something listening, not something recording.

As the user speaks, transcription appears on screen in real time in a large, readable text field. When they stop speaking, the transcription sits in an editable field they can correct, add to, or leave as is. This design removes the anxiety about imprecision in speech — they know they can fix it. So they speak freely.

A quiet "prefer to type instead?" affordance sits below the listening indicator. Small. Not drawing attention. Available if needed.

## Type default questions — these questions need precision

The following questions require exact information that transcription errors would corrupt:

- Products and services with specific names
- Prices
- Competitor names
- The signature phrases they always use — these need to be verbatim
- Words and phrases their content should never use
- Posting frequency preference
- Platform priorities

For these questions, the interface presents a clean text field with the cursor already active. No microphone prompt. No voice affordance drawing attention. Just the question and the space to answer.

A quiet "prefer to speak instead?" affordance is available below the text field for the one or two users who genuinely prefer it, but it is not the visual focus.

## **Hybrid questions — the user decides naturally**

Some questions work equally well in either mode. For these, the interface presents a voice listener and a text field side by side with equal visual weight. The user naturally gravitates to whichever feels right. No decision needs to be made — they either start speaking or start typing, and the interface responds accordingly.

The hybrid questions are:

- The caption the user would write tonight for their best product — this is their live voice sample
- What their customer's world looks like right now
- An example of content from another brand they admire

## **The one design mistake to avoid**

Do not present voice and type as a prominent binary choice at the start of every question. A toggle at the top of every screen asking "how would you like to answer?" creates decision fatigue and causes most users to default to typing — because typing feels safer. The result is a text-based onboarding with a voice feature nobody uses. All the infrastructure cost, none of the brand signal benefit.

The mode follows the question. The user feels like the interface knows what it's asking for. That is the right experience.

---

## Part Three: The Fourteen Questions — Full Specification

Below is the complete question architecture. For each question: the question text, the default input mode, what brand signal it is designed to extract, and any implementation notes.

---

### Question 1 — The plain language description

*"Describe what your business does in one or two sentences — the way you'd explain it to someone you just met, not the way you'd write it on a website."*

Default mode: Voice

Signal extracted: Natural vocabulary, genuine simplicity or complexity of the offering, how the owner frames their own business.

Implementation note: The phrase "not the way you'd write it on a website" is doing important work. It gives the user explicit permission to be casual. Do not remove it or soften it.

---

### Question 2 — The customer description

*"Who is your best customer? Don't describe them by age or income — describe them by what keeps them up at night, what they want more than anything, and why they came to you specifically."*

Default mode: Voice

Signal extracted: Customer psychology, emotional drivers, what the brand is actually solving at a human level.

Implementation note: The instruction "don't describe them by age or income" is critical. Without it, every user answers with demographics. Demographics are nearly useless for content generation. Anxieties and desires are the useful signal.

---

### Question 3 — Products and prices

*"What do you sell, and what does it actually cost? List your main products or services and their prices."*

Default mode: Type

Signal extracted: Product specifics, price anchors, portfolio breadth.

Implementation note: Allow multiple entries. Consider a simple repeating field — product name, price — that the user can add rows to. Keep it visually minimal. This is a data capture question, not a conversation question, and it should feel efficient.

---

#### **Question 4 — The live voice sample**

*"Write a caption for your best-selling product or service right now — the way you'd actually write it if you were posting it yourself tonight. Don't make it perfect. Just write it the way it naturally comes to you."*

Default mode: Hybrid — text field is prominent, voice is equally available

Signal extracted: This is the single most important question in the onboarding. The answer is the live voice sample — actual vocabulary, sentence rhythm, emoji usage, formality level, call-to-action style, cultural register. Everything the system needs to calibrate its generation layer.

Implementation note: This field must have no character counter, no formatting guidance, no example shown. The moment the user is given a format to follow, they follow it instead of writing naturally. Large open text area. No constraints. The instruction "don't make it perfect" is doing essential work — leave it exactly as written.

---

#### **Question 5 — Signature phrases**

*"What's a phrase or expression you use all the time when talking to customers — something that's just very you? Could be a greeting, a closing line, something you say when you're excited. Anything."*

Default mode: Type

Signal extracted: Signature vocabulary, recurring phrases, brand-specific expressions that become anchors in content generation.

Implementation note: Allow multiple entries. Users often have more than one. A simple "add another" affordance after the first entry is sufficient.

---

### **Question 6 — Aspirational voice reference**

*"Show me an example of content from another brand that you wish your content looked or sounded like. It can be a post, a website, anything. What specifically do you like about it?"*

Default mode: Hybrid

Signal extracted: Aspirational voice benchmark, specific stylistic preferences, gap between current voice and desired voice.

Implementation note: Allow either a URL paste or a text description of the content. Not everyone will have a URL ready. The text description of what they admire is often more useful than the link itself because it forces them to articulate what they're drawn to.

---

### **Question 7 — The point of view**

*"What's something your business genuinely believes that most of your competitors would never say publicly?"*

Default mode: Voice

Signal extracted: Brand conviction, competitive differentiation, intellectual positioning.

Implementation note: This question is intentionally slightly uncomfortable. Some users will pause before answering. Do not add helper text that softens the question or offers examples. The discomfort is productive — it's what produces honest, specific answers rather than safe, generic ones. Let the silence sit.

---

### **Question 8 — The recurring explanation**

*"What do your customers misunderstand about what you do or sell — the thing you find yourself explaining over and over again?"*



Default mode: Voice

Signal extracted: Educational content angles, common objections, the business owner's most persuasive natural language — because this is the explanation they've delivered many times and have refined.

Implementation note: Frame this as something they explain "over and over" rather than something customers "get wrong." The first framing positions the business owner as a patient educator. The second positions their customers as mistaken. The first produces warmer, more useful answers.

---

### **Question 9 — Source of pride**

*"What are you most proud of about your business — not what you're supposed to say, but what actually makes you feel like you're doing something worth doing?"*

Default mode: Voice

Signal extracted: Core brand values, authentic differentiators, emotional content anchors — the things the business genuinely cares about, which show in content that performs well.

Implementation note: The phrase "not what you're supposed to say" is essential. Remove it and users give the rehearsed answer. Keep it and many users give the honest answer, which is almost always more interesting and more useful.

---

### **Question 10 — Competitors**

*"Name two or three businesses you consider your main competitors — just the names that come to mind."*

Default mode: Type

Signal extracted: Competitive reference frame for avoidance and implicit differentiation.

Implementation note: Keep this simple. This is not a competitive analysis. You just need the names so content doesn't accidentally reference them favorably. Two to three is enough. Don't prompt for more.

---

## Question 11 — Customer's current world

*"What does your customer's world look like right now? What's happening in their life, their business, or their community that's relevant to what you sell?"*

Default mode: Hybrid

Signal extracted: Seasonal relevance, cultural moment awareness, customer life context — what makes content feel timely rather than generic.

Implementation note: This is the question that produces the most variance between users. Some will give a rich, specific answer that immediately informs the content strategy. Others will give a blank stare. If a user struggles with this, a single prompt — "Think about what time of year it is and what your customers are dealing with right now" — is sufficient to unlock an answer.

---

## Question 12 — Guardrails

*"What should your content never say or do? Any topics, tones, claims, or approaches that would feel wrong or damage your brand?"*

Default mode: Type

Signal extracted: Content guardrails, brand safety parameters, negative brand signals.

Implementation note: Negative constraints are as important as positive ones. A children's education brand that never wants political commentary, a health coach who never wants weight-loss before-and-after framing — these prevent the most trust-damaging content failures. Users answer this question more easily than most because violations feel immediately wrong to them.

---

## Question 13 — Posting parameters

*"How often do you want to post, and are there specific days or times when you know your audience is most active — even if it's just a feeling from experience?"*

Default mode: Type

Signal extracted: Posting frequency preference, timing intelligence, calendar parameters.

Implementation note: Offer sensible defaults that users can adjust — a simple frequency selector (daily, 5x/week, 3x/week, etc.) alongside a free-text field for timing notes. The defaults remove friction for users who don't have strong preferences. The free text captures the genuine knowledge of users who do.

---

### Question 14 — The brand personality

*"If your brand was a person at a dinner party, how would they behave? Would they be the loudest person in the room or the most interesting quiet one? Would they tell jokes or ask deep questions? Would they dress to impress or dress to be comfortable? Would they talk about themselves or ask about others?"*

Default mode: Voice

Signal extracted: Holistic personality profile, behavioral tendencies, emotional register — the gestalt of the brand personality in a form the AI can reason from.

Implementation note: This is the last question and the most open. The specific examples in the question — loudest or quietest, jokes or deep questions — are not exhaustive. They are prompts that unlock the user's imagination. The user will often go beyond the examples and describe their brand in ways that are unexpectedly precise and deeply useful. Let them run with it.

---

## Part Four: The Mid-Onboarding Wow Moment

This is a product feature, not a design flourish. It requires developer and designer to build it together deliberately.

**After questions 1 through 4 are complete** — once the system has the plain language description, the customer description, the product information, and the live voice sample — KratorAI generates a single piece of sample content in real time and shows it to the user before they continue.

Not at the end. Not as a preview of what's coming. Right there, in the middle of the onboarding flow.

The content shown should be one post — the type most appropriate for the user's primary platform — generated using the partial Brand Intelligence Profile built so far.

The moment it appears, the interface says something like: *"Here's what KratorAI already sounds like in your voice. We're only halfway through — it gets more accurate from here."*

### **Why this matters:**

Before this moment, the user is filling in an onboarding flow with uncertain faith that something useful will emerge. After this moment, they have seen evidence. They are now invested in completing the onboarding because they want to see how much better it gets with the remaining questions.

This single moment will measurably reduce onboarding abandonment. It converts a leap of faith into a demonstrated reality midway through the process.

### **Developer implementation note:**

This requires the generation layer to be callable with a partial Brand Intelligence Profile — not all fourteen questions answered, just the first four. The generation prompt for this mid-flow call should explicitly note that the profile is partial and instruct the model to produce its best output from available information. The result does not need to be perfect. It needs to be specific enough that the user recognizes their own brand in it.

---

## **Part Five: The Brand Intelligence Profile — Data Structure**

Everything captured in the fourteen questions gets structured into a Brand Intelligence Profile. This is the object your AI reasons from when generating every piece of content. It is not a static record — it is updated continuously throughout the user's time on the platform.

**Developer note:** The profile should be designed as a structured document that can be passed in full as context to the generation layer. Every field should be queryable individually for specific use cases (e.g., pulling only the guardrails for a safety check) and available in full for generation prompts.

---

### **BRAND IDENTITY**

- Business name

- Industry and category
  - Location and cultural context (city, country, cultural register)
  - Business stage: new / established / scaling
  - Price positioning: budget / mid / premium / luxury
- 

## OFFERING MAP

- Products and services — name, price, flagged as bestseller or seasonal
  - Products or services to never promote — discontinued, out of stock
  - Seasonal offerings with timing notes
- 

## AUDIENCE PROFILE

- Customer psychographic description (from Question 2 — stored verbatim, not summarized)
  - Core customer anxiety — what keeps them up at night
  - Core customer desire — what they want most
  - Why they chose this brand specifically
  - Customer's current moment — what's happening in their world right now
- 

## VOICE FINGERPRINT

- Formality level: 1–5 (1 = very casual / Lagos slang level, 5 = very formal)
- Warmth level: 1–5
- Humor presence: none / occasional / frequent
- Cultural register: specific expressions and references that are natural for this brand
- Emoji usage: none / minimal / moderate / heavy
- Caption length preference: short under 50 words / medium 50–150 / long over 150
- Signature phrases: verbatim list from Question 5, updated over time
- Never-use words and phrases: verbatim list from Question 12
- **Live voice sample: the verbatim caption from Question 4 — this is the most important field in the entire profile**
- Aspirational voice reference: URL or description from Question 6
- Running approved post samples: the last 3–5 posts the user approved without editing, stored verbatim — updated continuously

---

## **POINT OF VIEW**

- Core brand belief (from Question 7)
- Primary educational angle — the thing they explain repeatedly (from Question 8)
- Source of authentic pride (from Question 9)
- Brand personality archetype (from Question 14 — stored verbatim)

---

## **COMPETITIVE CONTEXT**

- Competitor names (from Question 10)
- Implicit differentiation notes — derived by the system from what the brand believes vs. what category competitors typically say

---

## **CONTENT PARAMETERS**

- Posting frequency
- Platform priorities with ranking
- Preferred posting times and days
- Content guardrails (from Question 12)
- Content pillars — these are derived by the system from the full profile, not asked directly

---

## **PERFORMANCE INTELLIGENCE (empty at onboarding, populated over time)**

- Best performing content types
  - Best performing posting times
  - Audience engagement patterns
  - Month-over-month strategic adjustments
  - Edit pattern analysis — what the user changes in generated content
-

## Part Six: The Confirmation Screen

After all fourteen questions are complete and before the calendar is generated, show the user a summary of what KratorAI learned about their brand.

This is not raw transcription. It is a clean, human-readable version of their Brand Intelligence Profile — written in plain language that a business owner can read and recognize themselves in.

Every section of the summary should be editable. A user who sees their brand described slightly wrong should be able to correct it immediately. Every correction is itself brand signal — it tells the system where its understanding needs calibration.

**Designer note:** The confirmation screen should feel like a reveal, not a review. The visual design should give the user the sense that the system genuinely learned something about them — not that they're proofing a database entry. The language of each section should be warm and specific, not clinical.

**Developer note:** Every edit made on this screen triggers an update to the corresponding field in the Brand Intelligence Profile. Edits made here are flagged as high-confidence corrections because the user made them deliberately after reviewing the system's understanding.

---

## Part Seven: The Calendar Reveal

This is not an end screen. It is a performance.

When the first 30-day content calendar is ready, present it as a reveal rather than a display.

Show the first week — not all 30 days at once — with each post showing: the day, the platform, the content type, and a preview of the caption. Captions are expandable. The user clicks to read the full post.

The moment a user reads the first full caption and recognizes their own voice in it is the second major wow moment of the onboarding. Design the screen around that moment. Give the caption enough visual space to be read properly. Do not bury it in a dense calendar grid.

Under the calendar, one line of copy does critical emotional work: *"KratorAI will publish these automatically. You don't need to do anything. But you can review, edit, or approve any post before it goes out."*

This removes obligation without removing control. Relief plus safety. That is the emotional state the user should leave onboarding in.

---

## Part Eight: The Edit-Capture System — How the Profile Gets Richer Over Time

This section is primarily for the developer and is one of the most important architectural decisions in the entire product.

Every time a user edits a generated post, they are teaching KratorAI something specific about their brand voice. The system must capture this signal automatically.

### What to capture:

When a user edits a post before approving it, the system stores both the original generated version and the edited version. It then compares the two and extracts the delta — what was added, what was removed, what was changed.

Over time, patterns in these deltas update the Voice Fingerprint automatically:

- User consistently removes exclamation marks → formality level increases
- User consistently adds a specific phrase → that phrase gets added to signature phrases
- User consistently removes a word → that word gets added to the never-use list
- User consistently shortens captions → length preference adjusts downward
- User consistently changes the call-to-action format → CTA pattern in the profile updates

### What to capture from approvals:

When a user approves a post without any edits, that post is added to the Running Approved Post Samples field in the Voice Fingerprint — stored verbatim. These become the most reliable voice examples in the profile because they represent content the user considered accurate enough to publish as-is.

When the Running Approved Post Samples field contains 5 or more posts, these should be weighted more heavily in the generation prompt than the original



onboarding voice sample, because they represent the system's most calibrated understanding of the brand voice to date.

### **What to capture from rejections:**

When a user rejects a post outright, the system flags the content type, angle, or structural pattern of that post as low-fit. Repeated rejections of the same content type update the guardrails automatically.

**Developer note:** This system does not require human review or manual tagging. It runs automatically in the background on every user interaction. The result over six months of use is a Brand Intelligence Profile that is dramatically richer and more calibrated than anything that could be captured at onboarding — built from hundreds of micro-corrections the user made naturally, without ever being aware they were training the system.

This is the compounding mechanism that makes KratorAI harder to leave over time. At month one, the voice is good. At month six, it approaches uncanny. A user who considers switching to a competitor at month six discovers that no competitor sounds like them, because no competitor has been listening and learning for six months.

---

## **Summary: What Developer and Designer Need to Build Together**

### **For the designer:**

The experience arc from anticipation through excitement governs every screen. Every design decision should be tested against the question: does this feel like a conversation or a form? The mid-onboarding wow moment and the calendar reveal are the two highest-stakes screens in the entire product — invest disproportionate design time in them. The voice/type split must feel seamless, with the mode following the question rather than the user being asked to choose.

### **For the developer:**

The Brand Intelligence Profile is a living document, not a static record. It must be structured to be passed as full context to the generation layer, updatable in real time from user interactions, and queryable by field for specific use cases. The generation layer must be callable with a partial profile for the mid-onboarding wow moment.

The edit-capture system runs automatically on every user interaction and is one of the most strategically important features in the product — build it from the start, not as a later addition.

**Together:**

The voice mode transcription, the mid-flow content generation, and the edit-capture system all require close coordination between design and engineering. Design the experience first — what should the user feel at each moment — then engineer to enable it. Do not let engineering constraints dictate the experience design without a conversation about what the constraint costs the user.

---

That is the full implementation brief. It covers the experience design, the question architecture with full rationale, the input system, the data structure, the wow moments, and the ongoing enrichment mechanism. Your team should be able to build from this directly.